

CS11-747 Neural Networks for NLP
Why is word2vec so fast?
Efficiency tricks for neural nets

Graham Neubig



Carnegie Mellon University
Language Technologies Institute

Site
<https://phontron.com/class/nn4nlp2019/>

Glamorous Life of an AI Scientist

Perception



Reality

```
neubig@itachi:~$ python nn-lm.py
[dynet] random seed: 3454201866
[dynet] allocating memory: 512MB
[dynet] memory allocation done.
---finished 500 sentences
---finished 1000 sentences
---finished 1500 sentences
---finished 2000 sentences
---finished 2500 sentences
---finished 3000 sentences
---finished 3500 sentences
---finished 4000 sentences
```

Waiting....

Why are Neural Networks Slow and What Can we Do?

- Big operations, especially for softmaxes over large vocabularies
 - → **Approximate operations or use GPUs**
- GPUs love big operations, but hate doing lots of them
 - → **Reduce the number of operations** through optimized implementations or batching
- Our networks are big, our data sets are big
 - → **Use parallelism** to process many data at once


Sampling-based Softmax Approximations

A Visual Example of the Softmax

$p = \text{softmax}(W h + b)$

Sampling-based Approximations

- Calculate the denominator over a subset



- Sample negative examples according to distribution q

Softmax

- Convert scores into probabilities by taking the exponent and normalizing (softmax)

$$P(x_i \mid \mathbf{h}_i) = \frac{e^{s(x_i \mid \mathbf{h}_i)}}{\sum_{\tilde{x}_i} e^{s(\tilde{x}_i \mid \mathbf{h}_i)}}$$

This is expensive, would like to approximate

$$Z(\mathbf{h}_i) = \sum_{\tilde{x}_i} e^{s(\tilde{x}_i \mid \mathbf{h}_i)}$$

Importance Sampling

(Bengio and Senecal 2003)

- Sampling is a way to approximate a distribution we cannot calculate exactly
- **Basic idea:** sample from arbitrary distribution Q (uniform/unigram), then re-weight with e^s/Q to approximate denominator

$$Z(\mathbf{h}_i) \approx \frac{1}{N} \sum_{\tilde{x}_i \sim Q(\cdot | \mathbf{h}_i)} \frac{e^{s(\tilde{x}_i | \mathbf{h}_i)}}{Q(\tilde{x}_i | \mathbf{h}_i)}$$

- This is a biased estimator (esp. when N is small)

Noise Contrastive Estimation

(Mnih & Teh 2012)

- **Basic idea:** Try to guess whether it is a true sample or one of N random noise samples. Prob. of true:

$$P(d = 1 \mid x_i, \mathbf{h}_i) = \frac{P(x_i \mid \mathbf{h}_i)}{P(x_i \mid \mathbf{h}_i) + N * Q(x_i \mid \mathbf{h}_i)}$$

- Optimize the probability of guessing correctly:

$$\mathbb{E}_P[\log P(d = 1 \mid x_i, \mathbf{h}_i)] + N * \mathbb{E}_Q[\log P(d = 0 \mid x_i, \mathbf{h}_i)]$$

- During training, approx. with unnormalized prob.

$$\tilde{P}(x_i \mid \mathbf{h}_i) = P(x_i \mid \mathbf{h}_i) / e^{c_{\mathbf{h}_i}} \quad (\text{set } c_{\mathbf{h}_i} = 0)$$

Simple Negative Sampling

(Mikolov 2012)

- Used in word2vec
- Basically, sample one positive k negative examples, calculate the log probabilities

$$P(d = 1 \mid x_i, \mathbf{h}_i) = \frac{P(x_i \mid \mathbf{h}_i)}{P(x_i \mid \mathbf{h}_i) + 1}$$

- Similar to NCE, but biased when $k \neq |V|$ or Q is not uniform

Mini-batching Negative Sampling

- Creating and arranging memory on the is expensive, especially on the GPU
- **Simple solution:** select the same negative samples for each minibatch
- (See Zoph et al. 2015 for details)

Let's Try it Out!

wordemb-negative-sampling.py

Structure-based Softmax Approximations

Structure-based Approximations

- We can also change the structure of the softmax to be more efficiently calculable
 - **Class-based softmax**
 - **Hierarchical softmax**
 - **Binary codes**

Class-based Softmax

(Goodman 2001)

- Assign each word to a class
- Predict class first, then word given class

$$P(c|h) = \text{softmax}(W_c h + b_c)$$


$$P(x|c,h) = \text{softmax}(W_x h + b_x)$$

- **Quiz:** What is the computational complexity?

Hierarchical Softmax

(Morin and Bengio 2005)

- Create a tree-structure where we make one decision at every node



- **Quiz:** What is the computational complexity?

Binary Code Prediction

(Dietterich and Bakiri 1995, Oda et al. 2017)

- Choose all bits in a single prediction

$$\sigma(W_c h + b_c) = \begin{matrix} 0 \\ 1 \\ 1 \\ 1 \\ 0 \\ \downarrow \\ \text{word 14} \end{matrix}$$




- Simpler to implement and fast on GPU

Let's Try it Out!



wordemb-binary-code.py

Two Improvement to Binary Code Prediction

Hybrid Model




Error Correcting Codes




Parallelism in Computation Graphs

Three Types of Parallelism

- Within-operation parallelism
 - Operation-wise parallelism
 - Example-wise parallelism
- 
- The diagram illustrates the three types of parallelism. It shows three bullet points on the left: 'Within-operation parallelism', 'Operation-wise parallelism', and 'Example-wise parallelism'. A large curly brace on the right groups the first two points under the heading 'Model parallelism'. Another curly brace further to the right groups all three points under the heading 'Data parallelism'.
- } Model parallelism
- } Data parallelism


Within-operation Parallelism



- GPUs excel at this!
- Libraries like MKL implement this on CPU, but gains less striking.
- Thread management overhead is counter-productive when operations small.

Operation-wise Parallelism


- Split each operation into a different thread, or different GPU device



- **Difficulty:** How do we minimize dependencies and memory movement?

Example-wise Parallelism

- Process each training example in a different thread or machine



- **Difficulty:** How do we accumulate gradients and keep parameters fresh across machines?

GPU Training Tricks

GPUs vs. CPUs

CPU, like a motorcycle



Quick to start, top speed
not shabby



GPU, like an airplane



Takes forever to get off the
ground, but super-fast
once flying

A Simple Example

- How long does a matrix-matrix multiply take?



Practically

- Use **CPU for profiling**, it's often fast (esp. in DyNet) and you can run many more experiments
- For **many applications, CPU is just as fast** or faster than GPU:
NLP analysis tasks with small or complicated data/networks
- You see **big gains on GPU when** you have:
 - Very big networks (or softmaxes with no approximation)
 - Do mini-batching
 - Optimize things properly

Speed Trick 1: Don't Repeat Operations

- Something that you can do once at the beginning of the sentence, don't do it for every word!

Bad

```
for x in words_in_sentence:  
    vals.append(W * c + x)
```

Good

```
W_c = W * c  
for x in words_in_sentence:  
    vals.append(W_c + x)
```

Speed Trick 2: Reduce # of Operations

- e.g. can you combine multiple matrix-vector multiplies into a single matrix-matrix multiply? Do so!

Bad

```
for x in words_in_sentence:  
    vals.append(W * x)  
val = dy.concatenate(vals)
```

Good

```
X = dy.concatenate_cols(words_in_sentence)  
val = W * X
```

- DyNet's auto-batching does this for you (sometimes)

Speed Trick 3: Reduce CPU-GPU Data Movement

- Try to **avoid memory moves** between CPU and GPU.
- When you do move memory, try to do it as early as possible (GPU operations are asynchronous)

Bad

```
for x in words_in_sentence:  
    # input data for x  
    # do processing
```

Good

```
# input data for whole sentence  
for x in words_in_sentence:  
    # do processing
```

What About Memory?

- Many GPUs only have up to 12GB, so **memory is a major issue**
- **Minimize unnecessary operations**, especially ones over big pieces of data
- If absolutely necessary, **use multiple GPUs** (but try to minimize memory movement)

Let's Try It!

slow-impl.py

Questions?